


How To Build a Many-Model Forecasting Solution on Snowflake

This guide walks through the complete **many-model pattern** on Snowflake: fit hundreds or thousands of independent models in parallel using Snowpark Container Services (SPCS), persist them to a stage, register them in Model Registry as a single model version, and run server-side batch inference via `run_batch`.

The accompanying **reference notebooks** ship in snowflakeR: `inst/notebooks/workspace_parallel_spcs_setup.ipynb` and `workspace_parallel_spcs_demo.ipynb`. The code snippets below follow that pattern. Lines marked **## ADAPT** are the parts you replace for your own data and models.

1. Architecture Overview

The many-model workflow has four phases:



Two dispatch modes are available:

Mode	Mechanism	Best for
Tasks	Snowflake Tasks + SPCS jobs; fixed batch of chunks dispatched upfront	Known workload; all chunks must complete
Queue	Hybrid Table queue; dynamic feeder adds chunks as workers finish	Time-boxed runs; streaming; scale-out/in

The “screen then refine” pattern: run a fast Tasks pass (e.g. stepwise ARIMA) over all SKUs, capture accuracy metrics, then rank the worst performers and feed them into a time-boxed Queue run that applies a more expensive model contest (ARIMA + ETS + TBATS with holdout validation). The time budget is spent where it matters most.

2. Prerequisites

Snowflake objects

- A **compute pool** with enough nodes (e.g. CPU_X64_L, 4 nodes).
- An **image repository** with the dosnowflake-worker Docker image pushed.
- A **warehouse** (XS is fine for orchestration SQL).
- The role used must have **USAGE** on the compute pool and **CREATE SERVICE / CREATE TASK** privileges.

Local R environment

```
install.packages(c("snowflakeR", "foreach", "iterators", "forecast",  
                  "ggplot2", "yaml"))
```

Authentication

Configure ~/.snowflake/connections.toml with a named profile using key-pair (JWT) auth:

```
[my_demo_jwt]  
account    = "MYORG-MYACCOUNT"  
user       = "MY_USER"  
authenticator = "SNOWFLAKE_JWT"  
private_key_path = "~/snowflake/keys/rsa_key.p8"  
warehouse  = "MY_WH"  
database    = "MY_DB"  
role        = "MY_ROLE"
```

3. Config and Bootstrap

The YAML config file

All environment-specific values live in a single YAML file. The reference implementation uses snowflaker_parallel_spcs_benchmark_ak_small.yaml:

```
## ADAPT: replace every value below with your account's objects  
context:  
  warehouse: MY_WAREHOUSE  
  database: MY_DATABASE  
  schema: SOURCE_DATA  
  
parallel_lab:  
  clean_room: false  
  database: MY_DATABASE  
  schemas:  
    source_data: SOURCE_DATA  
    config: CONFIG  
    models: MODELS  
  warehouse: MY_WAREHOUSE  
  compute_pool: MY_COMPUTE_POOL  
  image_uri: "myorg-myacct.registry.../my_db/config/images/dosnowflake-worker:latest"  
  dosnowflake_stage_name: DOSNOWFLAKE_STAGE  
  queue_table: DOSNOWFLAKE_QUEUE  
  instance_family: CPU_X64_L  
  
demo_forecast_n_skus: 2000  
demo_tasks_chunks_per_job: 4
```

```
demo_queue_n_workers: 4
demo_queue_chunks_per_job: 8
```

Loading config and connecting

```
source("snowflakeR/inst/notebooks/parallel_spcs_workflow.R")

cfg <- parallel_lab_load_config("path/to/my_config.yaml")
conn <- snowflakeR::sfr_connect(
  name = "my_demo_jwt",
  private_key_file = "~/snowflake/keys/rsa_key.p8"
)
```

Bootstrap DDL

`parallel_lab_setup()` runs bootstrap SQL that creates:

- 3 schemas: SOURCE_DATA, CONFIG, MODELS
- Internal stage DOSNOWFLAKE_STAGE with DIRECTORY = (ENABLE = TRUE)
- Hybrid Table DOSNOWFLAKE_QUEUE (for queue mode)
- Tables: TRAINING_RESULTS, TRAINING_METRICS, FORECAST_RESULTS
- View: MODEL_INDEX (parses RELATIVE_PATH from the stage directory table into RUN_ID + PARTITION_KEY)

```
parallel_lab_setup(conn, cfg, create_series = TRUE, n_units = 2000, n_days = 1095)
```

Set `create_series = TRUE` for synthetic data (generates 3 years of daily observations per SKU with weekly + annual seasonality, trend changepoints, and occasional outliers), or `FALSE` if your source table already exists.

ADAPT: the `SERIES_EVENTS` schema is (`UNIT_ID` VARCHAR, `OBS_DATE` DATE, `Y` FLOAT). Replace this with your own data table's DDL.

4. The foreach Expression

The worker expression is the R code that runs **inside each SPCS container** for every unit (SKU/store/sensor/etc.). It receives an injected `unit_data` data.frame and a `unit_id` variable.

Anatomy of a worker expression

```
## This runs inside the SPCS container for each unit_id.
## `unit_data` is automatically injected by doSnowflake - it contains the rows
## from SERIES_EVENTS where UNIT_ID == unit_id.

suppressPackageStartupMessages(library(forecast))

## ADAPT: convert your data to a ts object with the right frequency
## frequency = 7 for daily data (weekly cycle); use 12 for monthly, 365.25 for
## daily with annual cycle, etc.
ts_full <- ts(unit_data$Y, frequency = 7)

## ADAPT: fit your model here (any R model that produces forecasts)
model <- auto.arima(ts_full, stepwise = TRUE)
fc <- forecast(model, h = 30L)
acc <- forecast::accuracy(model)
```

```
## Contract: return a list with these fields
list(
  unit_id   = unit_id,
  model     = as.character(fc$method),
  forecast  = as.numeric(fc$mean),
  n_obs     = nrow(unit_data),
  model_obj = model,           # saved as .rds on stage
  aicc      = model$aicc,
  rmse      = acc[1, "RMSE"],
  mape      = holdout_mape,    # see holdout validation below
  mase      = acc[1, "MASE"]
)
```

Key contract: the returned list must include `unit_id` (the partition key) and `model_obj` (the fitted model object — serialized as `.rds` on the stage). Additional fields (`forecast`, metrics) are collected in the results list.

Holdout validation

For fair cross-method comparison, compute an **out-of-sample holdout MAPE** by reserving the last 30 days (or 1/4 of the series) as a test set:

```
n <- length(ts_full)
holdout_n <- min(30L, as.integer(floor(n / 4)))

holdout_mape <- tryCatch({
  ts_train <- window(ts_full, end = time(ts_full)[n - holdout_n])
  ts_test  <- window(ts_full, start = time(ts_full)[n - holdout_n + 1])
  m_ho     <- auto.arima(ts_train, stepwise = TRUE)
  fc_ho    <- forecast(m_ho, h = holdout_n)
  mean(abs((ts_test - fc_ho$mean) / ts_test)) * 100
}, error = function(e) NA_real_)
```

5. Tasks Mode (Fixed Batch)

Tasks mode dispatches all chunks upfront using a Snowflake Task DAG. Each task launches an SPCS job service that processes one chunk of `unit_ids`.

```
tasks_out <- parallel_lab_run_tasks_demo(conn, cfg,
  run          = TRUE,
  arima_stepwise = TRUE,
  forecast_h    = 30L,
  save_models   = TRUE
)
```

How it works

1. `registerDoSnowflake(conn, mode = "tasks", ...)` sets up the backend.
2. The `foreach` loop iterates over `unit_ids`; `doSnowflake` chunks them into `chunks_per_job` groups.
3. Each chunk becomes an SPCS job service that:
 - Bulk-reads its partition of `SERIES_EVENTS` via the `data_query` spec
 - Runs `mclapply` over the units within the container
 - Writes `.rds` model files to `@STAGE/models/<run_id>/<unit_id>.rds`
 - Writes result `.rds` back to `@STAGE/<job>/results/result_<chunk>.rds`
4. The orchestrator polls until all chunks complete, then collects results.

Chunk sizing

Balance parallelism vs overhead:

- **Too few chunks** (e.g. 1): all work in one container, no parallelism across nodes.
- **Too many chunks** (e.g. 2000, one per SKU): excessive job-launch overhead.
- **Sweet spot**: 4–20 chunks for 2000 units. Each container gets 100–500 SKUs and uses `mclapply` for intra-node parallelism.

6. Queue Mode (Time-Boxed Streaming)

Queue mode uses a Hybrid Table as a work queue. A **dynamic feeder** inserts rows as workers pull them, up to a time budget.

```
queue_out <- parallel_lab_run_queue_timeboxed(conn, cfg,
  run           = TRUE,
  time_limit_sec = tasks_out$elapsed,    # match Tasks wall time
  model_contest  = TRUE,                 # ARIMA + ETS + TBATS contest
  forecast_h     = tasks_out$params$forecast_h,
  n_chunks       = 40L,
  unit_ids       = worst_skus           # worst-first from Tasks
)
```

Dynamic feeder pattern

Unlike Tasks, the Queue does **not** dispatch all work upfront. The feeder loop:

1. Checks how many chunks are PENDING in the queue.
2. If slots are available and the time budget hasn't expired, inserts the next chunk.
3. Launches SPCS workers that claim rows via compare-and-swap (CAS): `UPDATE ... SET STATUS = 'RUNNING' WHERE STATUS = 'PENDING' LIMIT 1`.
4. When the feeder's time budget expires, it stops feeding but waits for in-flight chunks to finish.

This gives a natural “time-boxed” behavior: you spend exactly as much wall time as you budget, processing as many chunks as workers can handle.

Accuracy-ranked feeding

By ranking SKUs worst-first (highest holdout MAPE from the Tasks run), the time budget is spent where it has the greatest impact:

```
worst_skus <- parallel_lab_rank_skus_by_accuracy(
  run_results = tasks_out,
  metric      = "mape",
  descending  = TRUE
)
```

Model contest

When `model_contest = TRUE`, each SKU's worker expression runs a holdout- validated competition:

1. Split the daily series into training (~1065 days) and test (30 days).
2. Fit three candidates:
 - `auto.arima(stepwise = FALSE)` (exhaustive ARIMA search)
 - `ets()` (exponential smoothing)
 - `tbats()` (Trigonometric seasonality, Box-Cox, ARMA errors, Trend, Seasonal — handles multiple seasonal periods natively, e.g. weekly + annual cycles in daily data)
3. The candidate with the lowest test-set MAPE wins.

4. The winner is refit on the full series for the production model.

TBATS is particularly well-suited to this contest because the daily synthetic data has **two nested seasonal cycles** (weekly period = 7, annual period = 365.25) that ARIMA and ETS cannot capture simultaneously.

ADAPT: replace the three candidate model families with whatever makes sense for your domain (Prophet, XGBoost, LSTM, etc.).

7. Metrics and Comparison

Persisting metrics

After each run, save per-unit accuracy metrics to the `TRAINING_METRICS` table:

```
parallel_lab_save_metrics(conn, cfg, tasks_out)
parallel_lab_save_metrics(conn, cfg, queue_out)
```

This writes (`MODEL_RUN_ID`, `UNIT_ID`, `ARIMA_MODEL`, `AICC`, `RMSE`, `MAPE`, `MASE`, `N_OBS`) for every successfully fitted unit. The `MAPE` column stores the **holdout** (out-of-sample) MAPE — not the training MAPE.

Why holdout MAPE?

Training MAPE reflects how well the model fits historical data. Holdout MAPE reflects how well it **predicts unseen data**. When comparing a simple stepwise ARIMA (Tasks) against a model contest (Queue), training MAPE will look similar because both methods fit the full series. The difference only shows up in out-of-sample performance.

Plotting

Single-run histogram:

```
parallel_lab_plot_accuracy(tasks_out, metric = "mape")
```

Faceted comparison (Tasks vs Queue, matched SKUs only):

```
parallel_lab_plot_accuracy_comparison(
  tasks_out, queue_out,
  metric = "mape",
  matched_only = TRUE
)
```

The `matched_only = TRUE` parameter restricts the comparison to SKUs that appear in both runs, making the distributions directly comparable.

8. Model Registry

Registering N models as one version

After training, each SKU's `.rds` model file lives at `@STAGE/models/<run_id>/<unit_id>.rds`. The stage's `DIRECTORY` table tracks all files. The `MODEL_INDEX` view extracts `RUN_ID` and `PARTITION_KEY` from the file paths.

`sfr_log_many_model()` registers a **CustomModel aggregator** in Model Registry. The aggregator:

- Carries a `model_index.json` mapping `UNIT_ID` → stage path
- Exposes `@partitioned_api predict()` for `run_batch` (`TABLE_FUNCTION`)
- Exposes `@inference_api predict_single()` for `SPCS` service (`FUNCTION`)
- Registers in $O(1)$ time (~8–13s regardless of partition count)

```
reg_tasks <- parallel_lab_register_models(conn, cfg,
  model_run_id = tasks_out$model_run_id,
  forecast_h   = tasks_out$params$forecast_h
)
```

The Ray shuffle bug (SNOW-2193288)

Model Registry's `run_batch` uses Ray internally for partitioned dispatch. When the number of unique partition keys is much less than ~200, Ray's shuffle mechanism can fail. **Always ensure ≥ 200 partition keys** when using `run_batch`. The reference implementation warns if this threshold is not met.

9. Batch Inference

Server-side partitioned inference

`sfr_run_batch()` launches an SPCS batch job that:

1. Partitions the input DataFrame by `UNIT_ID`.
2. For each partition, the `RManyModelAggregator` downloads the `.rds` model from stage, loads it via `rpy2`, and runs `forecast::forecast(model, h=...)`.
3. Returns a DataFrame of (`PARTITION_KEY`, `PERIOD`, `FORECAST`, `LOWER_80`, `UPPER_80`, `LOWER_95`, `UPPER_95`, `STATUS`).

```
fc_df <- parallel_lab_run_inference(conn, cfg,
  model_run_id = tasks_out$model_run_id,
  forecast_h   = tasks_out$params$forecast_h
)
```

The results are written to the `FORECAST_RESULTS` table with columns: `RUN_ID`, `UNIT_ID`, `HORIZON`, `FORECAST_DATE`, `POINT_FORECAST`, `LO_80`, `HI_80`, `LO_95`, `HI_95`, `STATUS`.

Input contract

The input DataFrame needs at minimum a `UNIT_ID` column (the partition key). Additional columns are ignored by the aggregator — it loads the model from stage rather than re-training.

10. Monitoring

The reference implementation includes two Streamlit dashboards:

- `streamlit_benchmark_monitor.py`: live progress for Tasks vs Queue benchmark runs, filtered by `JOB_ID`.
- `streamlit_parallel_demo_monitor.py`: general-purpose SPCS service and queue monitoring.

Deploy to Snowflake (Streamlit-in-Snowflake) with:

```
python deploy_streamlit_monitor.py \
  --connection-name my_demo_jwt \
  --private-key-file ~/.snowflake/keys/rsa_key.p8
```

Tip: scope dashboard queries by `JOB_ID` to avoid showing cumulative counts from previous runs.

11. Lessons Learned

JWT expiry and stale sessions

`sfr_connect()` auto-detects Workspace Notebook sessions. When running in RStudio with explicit `name` / `account` parameters, it skips auto-detection to avoid wrapping a stale, expired Python session. If you see `JWT token is invalid` after a timeout, explicitly reconnect:

```
conn <- snowflakeR::sfr_connect(
  name = "my_demo_jwt",
  private_key_file = "~/snowflake/keys/rsa_key.p8"
)
dbi_con <- snowflakeR::sfr_dbi_connection(conn)
```

RStudio Connections Pane

Call `sfr_dbi_connection(conn)` after `sfr_connect()` to populate the Connections Pane with database catalog info. If schemas or tables don't appear, disconnect in the pane and reconnect — RStudio caches connection observer closures.

Setting the right ts frequency

Use `ts(data, frequency = 7)` for daily data with a weekly cycle, or `frequency = 12` for monthly data. Without the correct frequency, `auto.arima` cannot learn seasonal patterns and forecasts will be flat. For daily data with both weekly and annual cycles, TBATS auto-detects multiple seasonal periods from the `frequency` attribute.

Chunk count parity

When comparing Tasks vs Queue, use chunk counts that give comparable total work. The Queue's dynamic feeder means fewer chunks may complete within the time budget — that's by design (time-boxed, not work-boxed).

bquote vs quote for parameter substitution

The worker expression uses `bquote()` with `.()` splicing to inject parameter values (e.g. `.(forecast_h)`, `.(arima_stepwise)`) at definition time. Using `quote()` would capture the variable *names*, not their *values*, causing failures when the expression runs in a remote container where those variables don't exist.

```
forecast_expr <- bquote({
  model <- auto.arima(ts_full, stepwise = .(arima_stepwise))
  fc    <- forecast(model, h = .(forecast_h))
  # ...
})
```

Further Reading

- **vignette("parallel-dosnowflake")**: the low-level `registerDoSnowflake` API reference (foreach backend, generic SPCS executor, crew/mirai workers).
- **Reference notebooks**: `inst/notebooks/workspace_parallel_spcs_setup.ipynb`, `workspace_parallel_spcs_demo.ipynb` contains the complete runnable demo with configs, deploy scripts, and tests.
- **Function reference**: `?sfr_log_many_model`, `?sfr_run_batch`, `?sfr_model_registry`, `?registerDoSnowflake`.